


初探 Drasi

微軟開源的 Data Change Processing Platform





Alan Tsai

蔡孟玔

- 後端工程師 - .NET 技術為主
- 喜歡學習不同東西
 - Azure
 - Data Science、Chatbot
 - Container、DevOps
 - 加強開發的Tools、架構
- 翻譯文章/軟體
 - 兩本翻譯書
- 三門線上課程
- 看小說、玩手游

教就是最好的學習方式 - 喜歡考證驗證所學



喜歡技術分享

.NET Conf
TAIWAN
2024

2024 11 月 GDG 小聚

探索 GitHub Copilot 的新功能讓我們的應用程式開發更快一步

2024 08 月 TSMC IT Community Meetup

聊聊生成式 AI 衝擊下工程師可以做好什麼準備

2024 06 月 Global DevOps Experience Day

2023 .NET Conf Taiwan

探索 API 開發的挑戰與解決之道

從 GitHub Copilot 到 Enterprise Copilot: 打造符合企業需求的智能開發助手之路

2023 10 月 GDG 小聚

開始邁上你的 Azure 成本管控大師之路

2022 .NET Conf Taiwan

談 Event Driven Architecture 之前是不是該把 Event 規格搞定?

CloudEvents 是什麼? | 邁上 Cloud Native App 之路

STUDY4.TW

為 學 習 而 生



喜歡技術分享

.NET Conf
TAIWAN
2024

Tibame

AZ-900、AZ-104、AZ-204

GitHub Copilot

企業導入，在製造業、金融業、服務業，已培訓近 2000+ 人次

Trainocate

AZ-104、AZ-204、AZ-400、DP-200、DP-201、PL-900、DA-100

台灣智慧自動化與機器人協會

運用Python進行大數據分析、機器學習基礎理論課程及人工智慧 – ML.NET

中華電信學院

DevOps與CI/CD實務研習班

Windows Container 技術實務班

使用Azure AI打造有人工智能的Line聊天機器人

緯育 TibaMe

TRAINOCATE

TAIROA
社團法人台灣智慧自動化與機器人協會

中華電信學院

教就是最好的學習方式 - 喜歡考證驗證所學

Alan Tsai 的學習筆記

<https://blog.alantsai.net>



contact@alantsai.net

FB 粉絲頁

<http://fb.alantsai.net>



contact@alantsai.net

影片

<http://yt.alantsai.net>



<http://bili.alantsai.net>



Agenda

01

What is the problem

02

What is Drasi

03


How to use Drasi

What is the problem

What is a Program

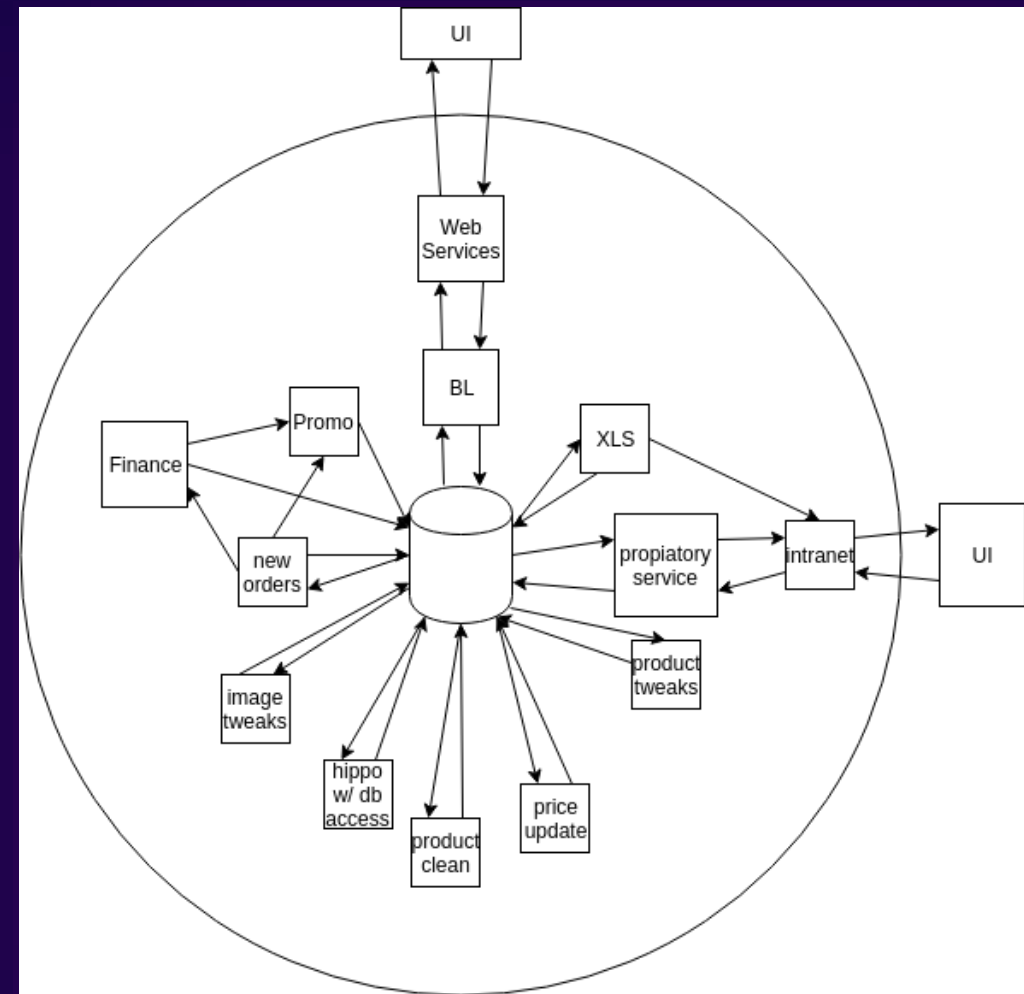

Data + Logic

Death Star Architecture



It appears
powerful but is
fundamentally
vulnerable.

1. Take the cheaper faster option.
2. Take the cheaper faster option.
3. Goto 1.



What is Death Star Architecture and why you should care

改 A 壞 B => 有程式碼品質問題

Gate Keepers 變成英雄

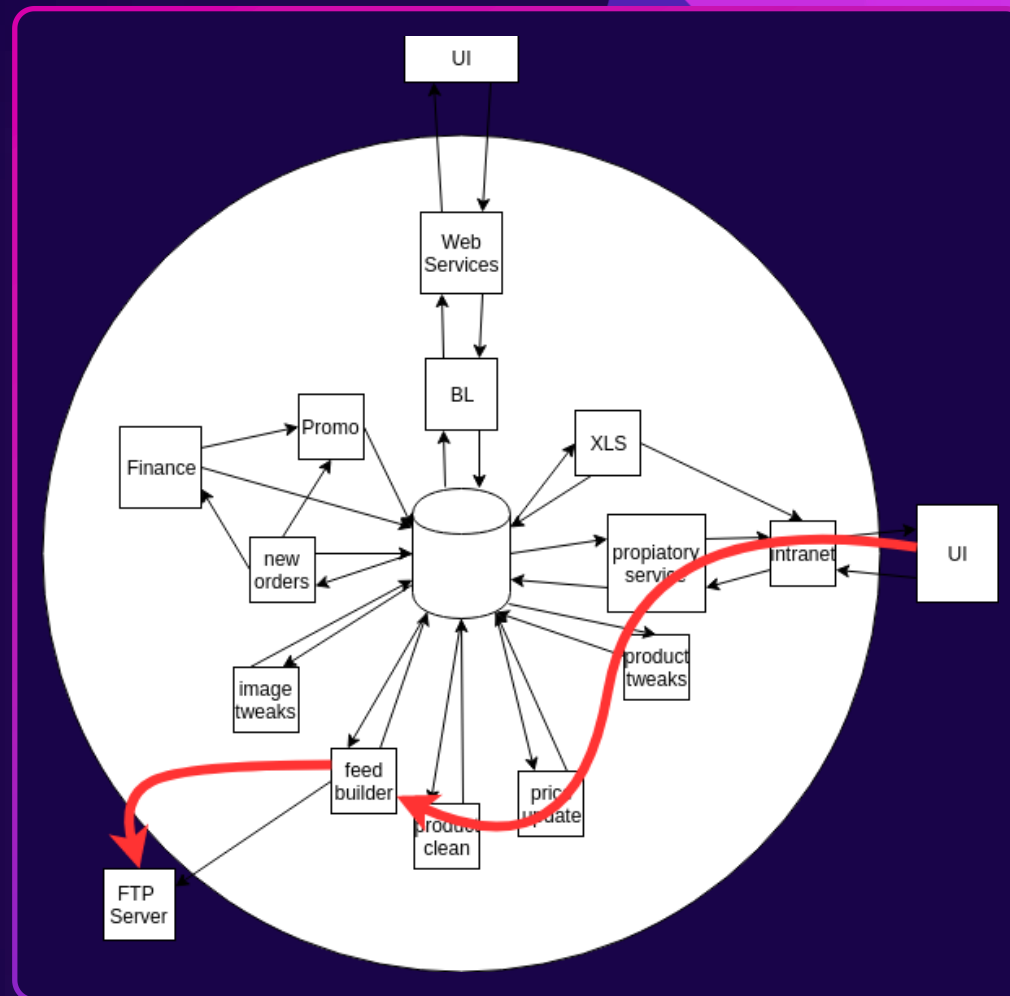
為了加速開發，我們導入“敏捷”專注在 Process, Documentation, Planning

整個糾結在一起，value stream 散落在各個解決方案

開始有很多分支 (但是 DB 是同一座...)

商業需求不會因為開發變慢而停下

各個部門解決自己的問題，而沒有思考整體公司。沒有一致步調前進



**We need to share Data
not Data Source**

Three way of communication Logic and Data

01

Synchronous
API

02

Asynchronous
Event

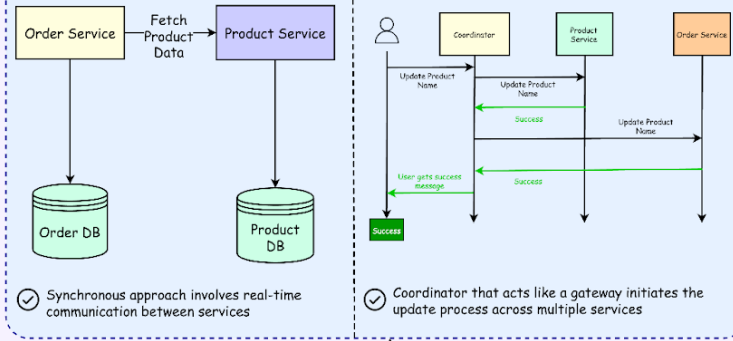
03

Data
CDC


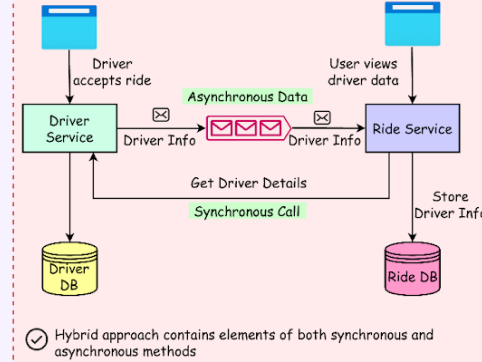
Organizations often need to take some action when data stored in a database changes, or, conversely, when it doesn't change as expected

A Cheatsheet On Data Sharing Between Microservices

Synchronous Data Sharing



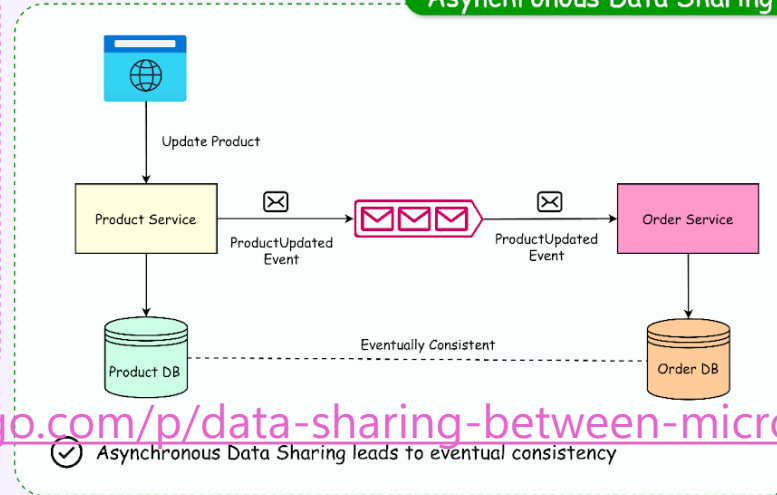
Hybrid Approach



Trade-offs and Decisions

- ✓ Evaluate the consistency requirements
- ✓ Strong vs Eventual Consistency
- ✓ Performance considerations such as latency, throughput, and resource utilization
- ✓ Application scalability considerations such as horizontal scaling, data volume, and service evolution

Asynchronous Data Sharing



<https://blog.bytebytego.com/p/data-sharing-between-microservices>

What is the challenge?



**Change Detection
Challenges**



**Real-Time Response
Issues**




**Developer and
Business Impact**

instantaneous query



What is the current solution



What is drasi

Propose to CNCF Sandbox Open Source Project



Change logs

Detect



Sources



Continuous Queries



Reactions

React

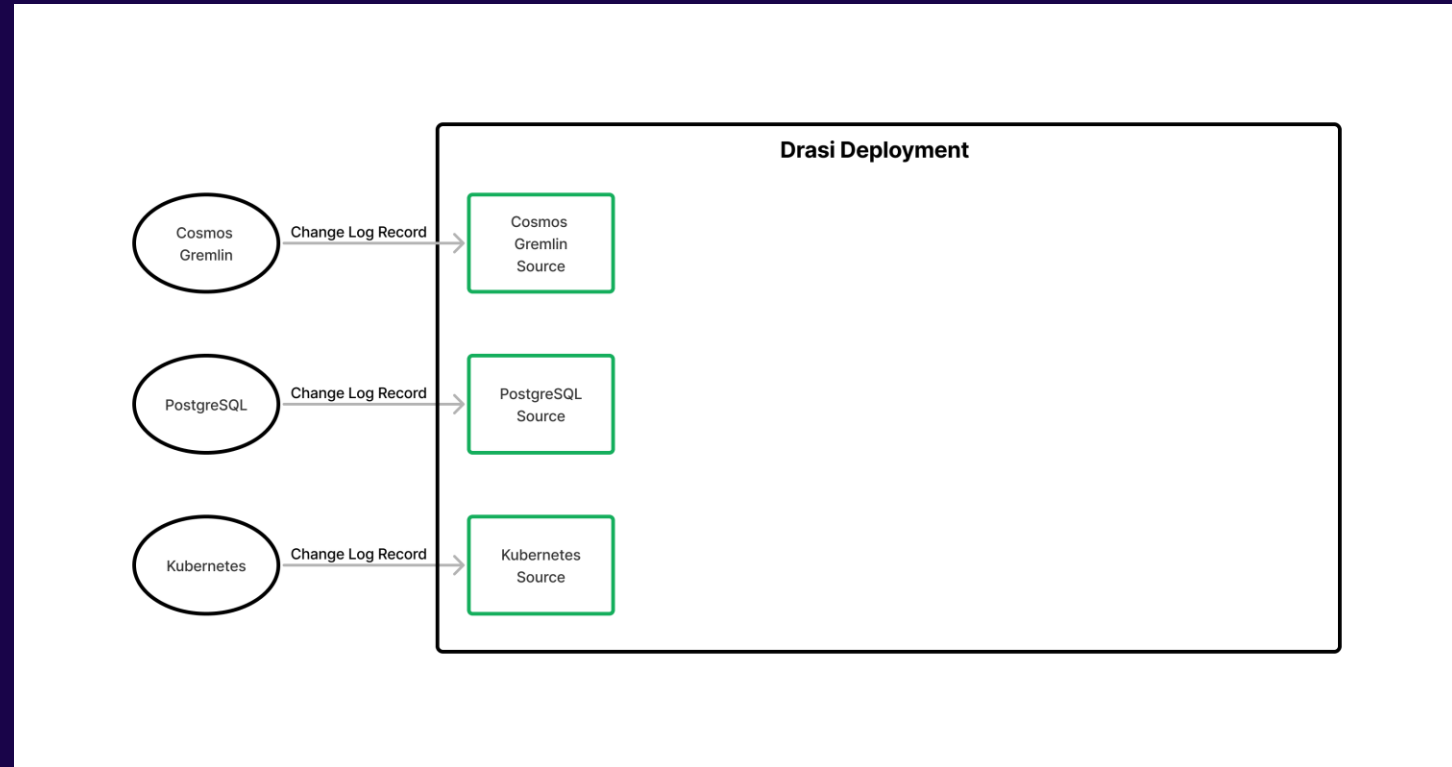


Action

Source

Connection of data

- [Azure Cosmos DB Gremlin API](#)
- [PostgreSQL](#)
- Event Hubs
- Microsoft SQL Server
- Microsoft Dataverse
- Kubernetes



Continuous Queries

Cypher query

- **Real-Time Data Updates:**


- Continuously incorporate changes from the source database.
- Maintain accurate, up-to-date query results.

- **Flexible Query Capabilities:**

- Use Cypher Query Language for detailed change detection.
- Rich query logic for properties and relationships.

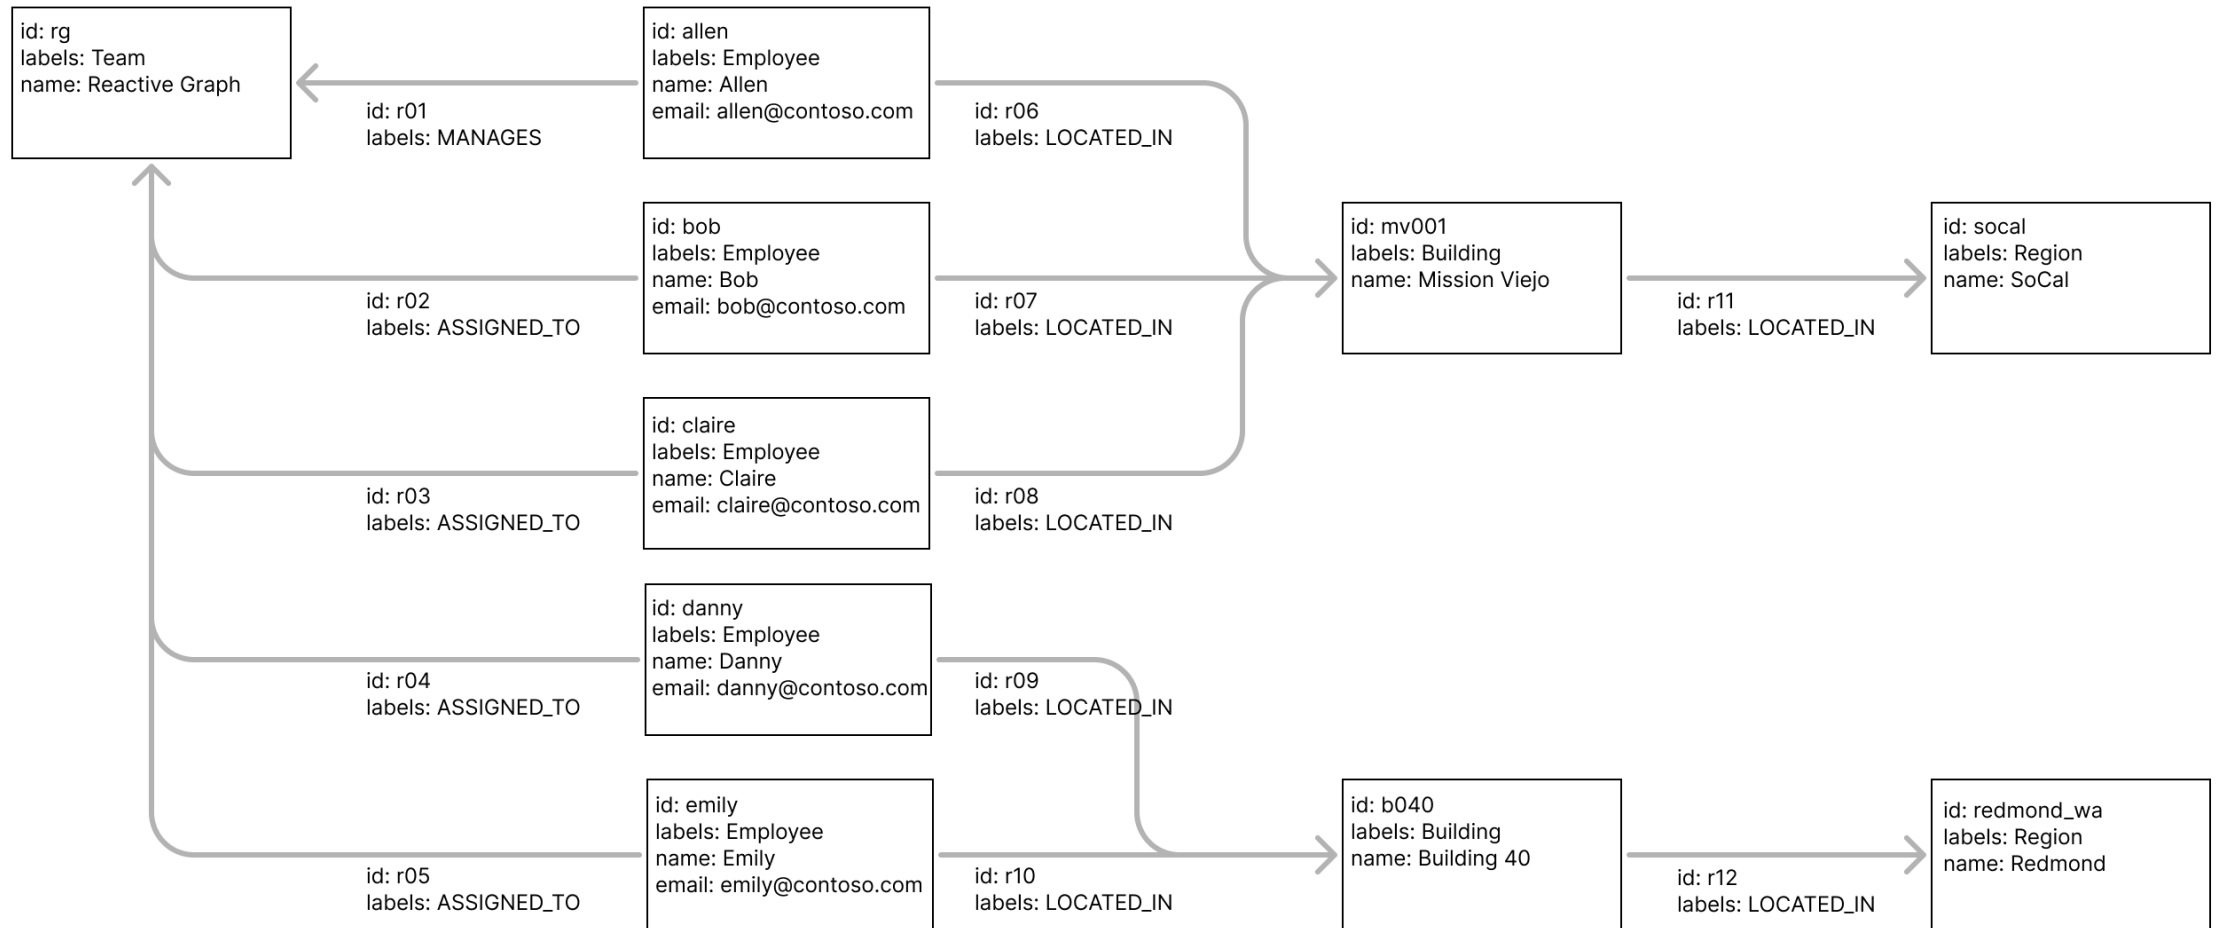
- **Multi-Source Integration:**

- Queries span multiple data sources without complex joins.
- Integrates relational and graph data seamlessly.

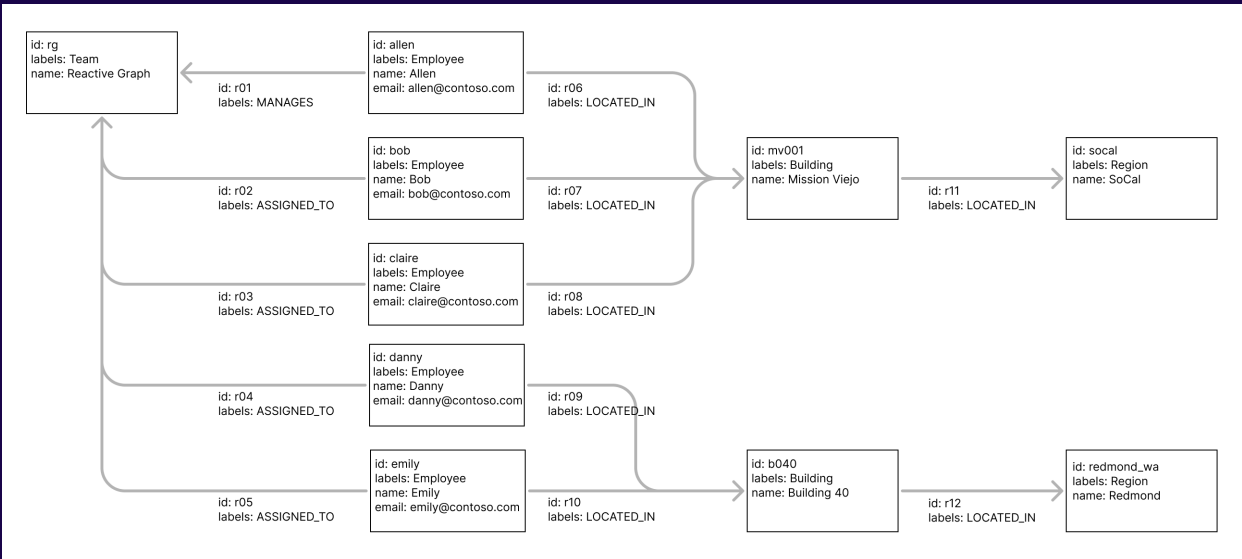


Incident Alert Detection

Cypher query: Node and Relation



Incident Alert Detection



Employees located in **Buildings** within **Regions** where there are active **Incidents** of **type** 'environmental' that have a **severity** level of 'critical' or 'extreme'

The name and email address of the at risk employee and their manager

MATCH

```

(e:Employee)-[:ASSIGNED_TO]->(t:Team),
(m:Employee)-[:MANAGES]->(t:Team),
(e:Employee)-[:LOCATED_IN]->(b:Building)-[:LOCATED_IN]->(r:Region),
(i:Incident {type:'environmental'})-[:OCCURS_IN]->(r:Region)

```

WHERE

```

elementId(e) <> elementId(m) AND i.severity IN ['critical', 'extreme'] AND i.endTimeMs IS NULL

```

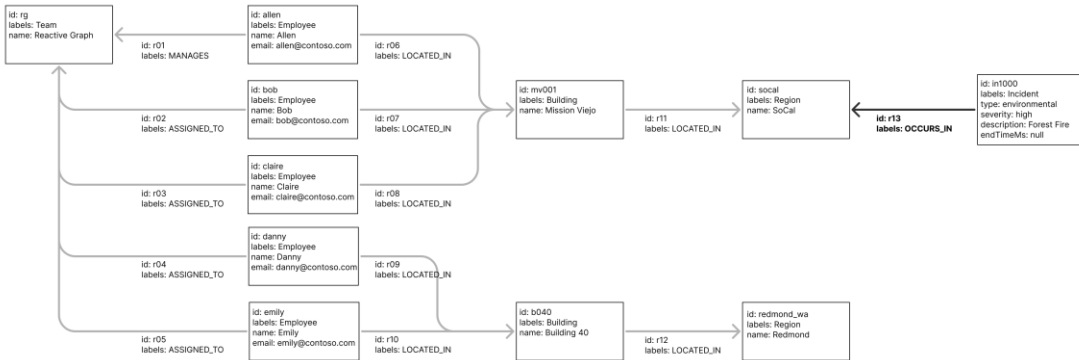
RETURN

```

m.name AS ManagerName, m.email AS ManagerEmail,
e.name AS EmployeeName, e.email AS EmployeeEmail,
r.name AS RegionName,
elementId(i) AS IncidentId, i.severity AS IncidentSeverity, i.description AS IncidentDescription

```

Incident Alert Detection



```
{
  "added": [
    { "ManagerName": "Allen", "ManagerEmail": "allen@contoso.com",
      "Southern California", "IncidentId": "in1000", "IncidentSeverity": "high",
      "IncidentDescription": "Forest Fire", "IncidentEndTimeMs": null },
    { "ManagerName": "Allen", "ManagerEmail": "allen@contoso.com",
      "Southern California", "IncidentId": "in1000", "IncidentSeverity": "high",
      "IncidentDescription": "Forest Fire", "IncidentEndTimeMs": null }
  ],
  "updated": [],
  "deleted": []
}
```

MATCH

```
(e:Employee)-[:ASSIGNED_TO]->(t:Team),
(m:Employee)-[:MANAGES]->(t:Team),
(e:Employee)-[:LOCATED_IN]->(b:Building)-[:LOCATED_IN]->(r:Region),
(i:Incident {type:'environmental'})-[:OCCURS_IN]->(r:Region)
```

WHERE

```
elementId(e) <> elementId(m) AND i.severity IN ['critical', 'extreme'] AND i.endTimeMs IS NULL
```

RETURN

```
m.name AS ManagerName, m.email AS ManagerEmail,
e.name AS EmployeeName, e.email AS EmployeeEmail,
r.name AS RegionName,
elementId(i) AS IncidentId, i.severity AS IncidentSeverity, i.description AS IncidentDescription
```

Incident Alert Detection

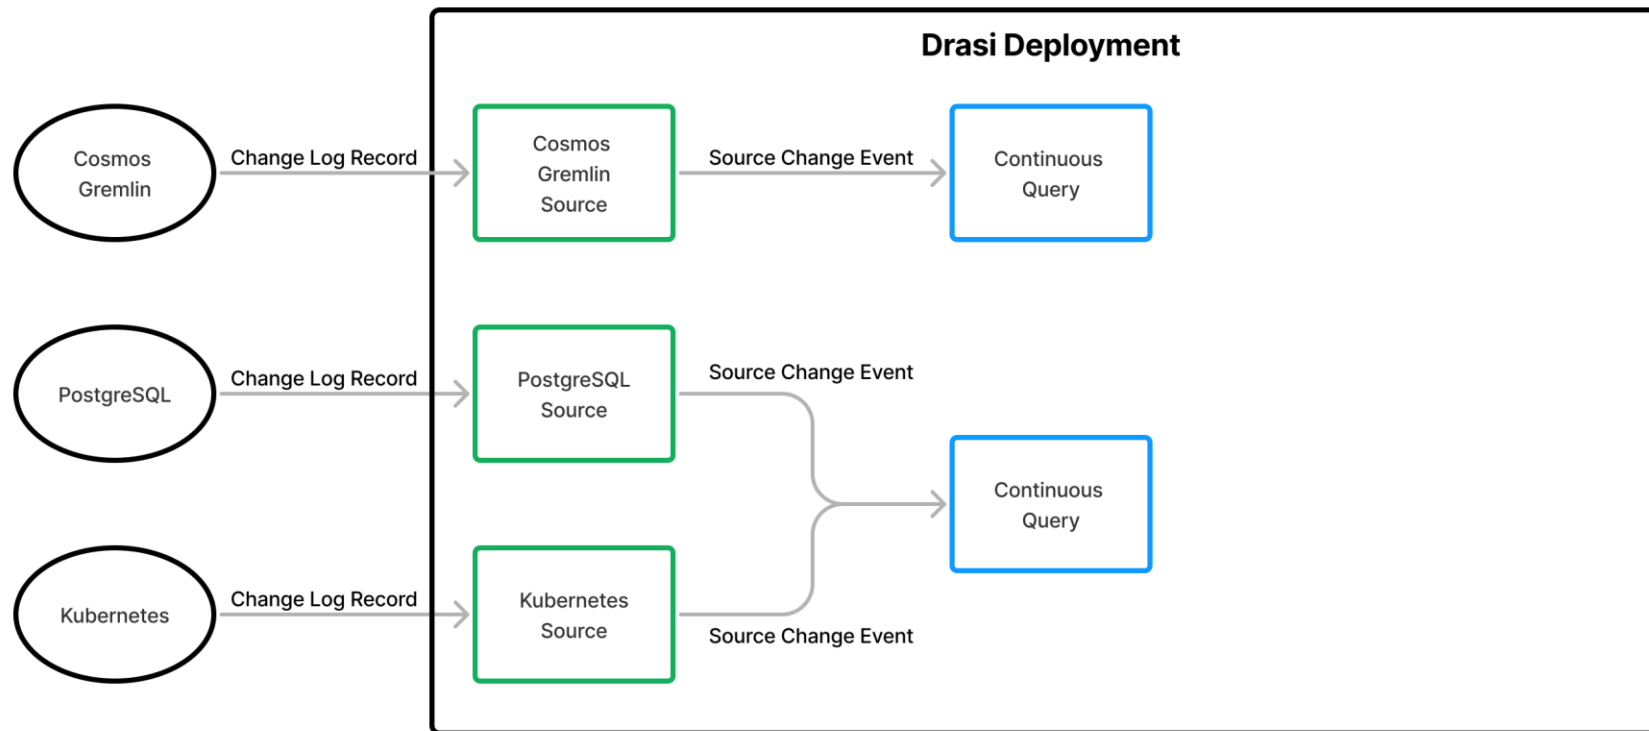
If Bob Change Location

```
{  
  "added": [],  
  "updated": [],  
  "deleted": [ {  
    "ManagerName": "Allen",  
    "ManagerEmail":  
    "allen@contoso.com",  
    "EmployeeName": "Bob..." }  
  ]  
}
```

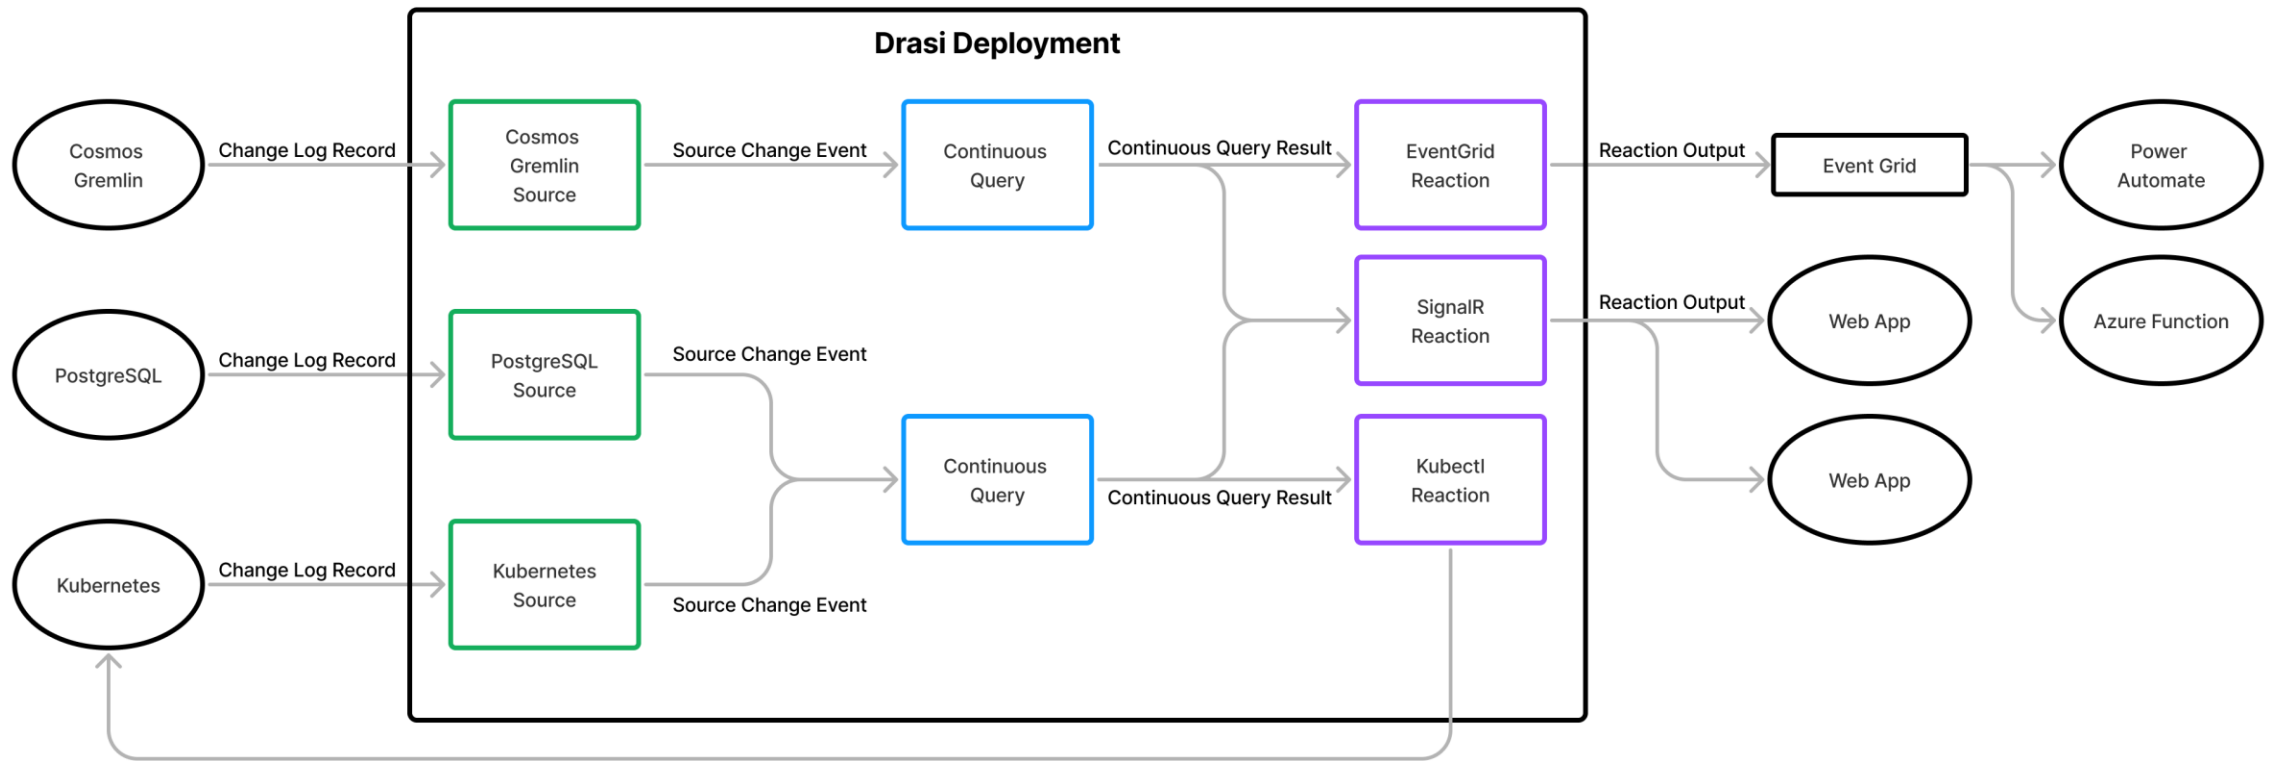
If severity change

```
{  
  "added": [],  
  "updated":  
  [ { "before": { "ManagerName": "Allen",  
    "ManagerEmail": "allen@contoso.com",  
    "EmployeeName": "Bob..." },  
    "after": { "ManagerName": "Allen",  
    "ManagerEmail": "allen@contoso.com",  
    "EmployeeName": "Bob..." } } ],  
  "deleted": []  
}
```

Continuous Query Abstract



Reaction



How to use drasi

Propose to CNCF Sandbox Open Source Project

drasi cli

Dependency

1. Kubernetes
2. Dapr



Monitor

1. OpenTelemetry
2. Prometheus



Install on AKS

1. Connect to AKS

```
az aks get-credentials  
--resource-group <group>  
--name <name>
```

2. Download drasi cli

```
curl -fsSL  
https://raw.githubusercontent.com/drasi-  
project/drasi-  
platform/main/cli/installers/install-drasi-  
cli.sh | /bin/bash
```

3. Call drasi init

```
alan [ ~ ]$ drasi init  
Installing Drasi with version 0.1.6 from registry ghcr.io  
✓ Dapr installed successfully  
✓ Infrastructure deployed  
  ✓ app=drasi-redis is online  
  ✓ app=drasi-mongo is online  
✓ Control plane is online  
  ✓ drasi/infra=api is online  
  ✓ drasi/infra=resource-provider is online  
✓ Query container created  
  ✓ Apply: QueryContainer/default: complete  
  ✓ Wait QueryContainer/default online  
✓ Default source providers created  
  ✓ Apply: SourceProvider/PostgreSQL: complete  
  ✓ Apply: SourceProvider/SQLServer: complete  
  ✓ Apply: SourceProvider/CosmosGremlin: complete  
  ✓ Apply: SourceProvider/Dataaverse: complete  
  ✓ Apply: SourceProvider/EventHub: complete  
✓ Default reaction providers created  
  ✓ Apply: ReactionProvider/Debug: complete  
  ✓ Apply: ReactionProvider/Debezium: complete  
  ✓ Apply: ReactionProvider/EventGrid: complete  
  ✓ Apply: ReactionProvider/Gremlin: complete  
  ✓ Apply: ReactionProvider/Result: complete  
  ✓ Apply: ReactionProvider/SignalR: complete  
  ✓ Apply: ReactionProvider/StorageQueue: complete  
  ✓ Apply: ReactionProvider/StoredProc: complete  
  ✓ Apply: ReactionProvider/Dataaverse: complete
```

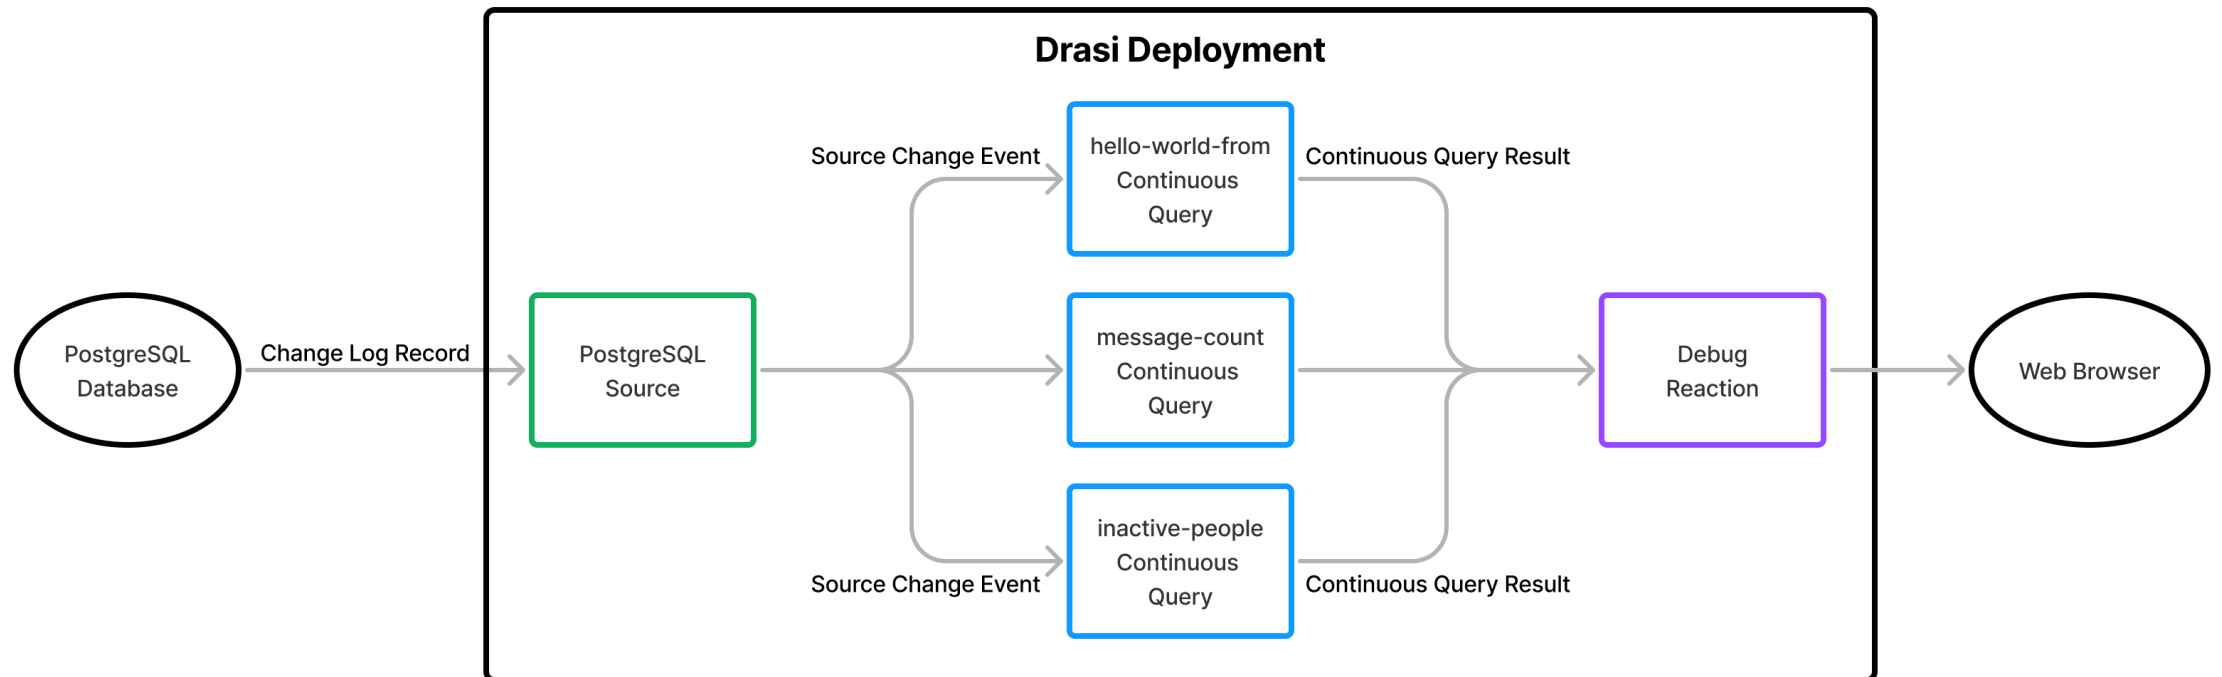
```
alan [ ~ ]$ kubectl get pods -n dapr-system
```

NAME	READY	STATUS	RESTARTS	AGE
dapr-dashboard-5cc65d985f-zkbsc	1/1	Running	0	6m52s
dapr-operator-5d98f57c86-pcjft	1/1	Running	1 (6m28s ago)	6m52s
dapr-placement-server-0	1/1	Running	0	6m52s
dapr-sentry-697bdc6cc4-g2gdj	1/1	Running	0	6m52s
dapr-sidecar-injector-56c4c4b485-25gvp	1/1	Running	0	6m52s

<https://drasi.io/how-to-guides/installation/install-on->

Solution Architecture

1. Which people have sent the message "Hello World"?
2. How many times has the same message been sent?
3. Which people haven't sent a message in the last 20 seconds?



Prepare the data - PostgreSQL

Field Name	Type	Description
MessageId	integer	A unique id for each message.
From	character varying(50)	The name of who the message is from.
Message	character varying(200)	The text of the message.

MessageId	From	Message
1	Buzz Lightyear	To infinity and beyond!
2	Brian Kernighan	Hello World
3	Antoninus	I am Spartacus
4	David	I am Spartacus

Deploy PostgreSQL

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: postgres
spec:
  replicas: 1
  selector:
    matchLabels:
      app: postgres
  template:
    metadata:
      labels:
        app: postgres
    spec:
      containers:
        - name: postgres
          image: postgres:15-alpine
          args: ["-c", "wal_level=logical"]
          volumeMounts:
            - name: init
              mountPath: "/docker-entrypoint-initdb.d"
          ports:
            - containerPort: 5432
          envFrom:
            - configMapRef:
                name: test-pg-config
      volumes:
        - name: init
          configMap:
            name: test-data-init
```

```
kubectl apply
-f drasi-postgres.yaml
```

```
kubectl port-forward
svc/postgres
```

5432:5432

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: test-data-init
data:
  init.sql: >
    CREATE TABLE "Message" (
      "MessageId" integer NOT NULL,
      "From" character varying(50) NOT NULL,
      "Message" character varying(200) NOT NULL
    );

    ALTER TABLE "Message" ADD CONSTRAINT pk_message
    PRIMARY KEY ("MessageId");

    INSERT INTO public."Message" VALUES (1, 'Buzz Lightyear', 'To infinity and beyond!');
    INSERT INTO public."Message" VALUES (2, 'Brian Kernighan', 'Hello World');
    INSERT INTO public."Message" VALUES (3, 'Antoninus', 'I am Spartacus');
    INSERT INTO public."Message" VALUES (4, 'David', 'I am Spartacus');
```

```
apiVersion: v1
kind: Service
metadata:
  name: postgres
  labels:
    app: postgres
spec:
  type: ClusterIP
  ports:
    - port: 5432
  selector:
    app: postgres
```

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: test-pg-config
  labels:
    app: postgres
data:
  POSTGRES_DB: hello-world
  POSTGRES_USER: test
  POSTGRES_PASSWORD: test
```

Test PostgreSQL

```
psql -h localhost -U test  
-d hello-world
```

```
alan [ ~ ]$ psql -h localhost -U test -d hello-world  
psql (14.13, server 15.10)  
WARNING: psql major version 14, server major version 15.  
Some psql features might not work.  
Type "help" for help.
```

```
hello-world=# select * from public.message  
hello-world=# SELECT * FROM "Message";  
ERROR: syntax error at or near "SELECT"  
LINE 2: SELECT * FROM "Message";
```

```
^  
hello-world=# ^C  
SELECT * FROM "Message";  
hello-world=# ^C  
hello-world=# ^C  
hello-world=# ^C
```

```
hello-world=# SELECT * FROM "Message";  
 MessageId |      From      |      Message  
-----+-----+-----  
          1 | Buzz Lightyear | To infinity and beyond!  
          2 | Brian Kernighan | Hello World  
          3 | Antoninus      | I am Spartacus  
          4 | David          | I am Spartacus  
(4 rows)
```

```
hello-world=#
```

Source

```
apiVersion: v1
kind: Source
name: hello-world
spec:
  kind: PostgreSQL
  properties:
    host: postgres.default.svc.cluster.local
    port: 5432
    user: test
    password: test
    database: hello-world
    ssl: false
    tables:
      - public.Message
```

```
alan [ ~ ]$ ./drasi apply -f ./hello-world-source.yaml
✓ Apply: Source/hello-world: complete
alan [ ~ ]$ drasi list source
bash: drasi: command not found
alan [ ~ ]$ ./drasi list source
      ID      | AVAILABLE | MESSAGES
-----+-----+-----
  hello-world | true      |
alan [ ~ ]$
```

Query

```
apiVersion: v1
kind: ContinuousQuery
name: hello-world-from
spec:
  mode: query
  sources:
    subscriptions:
      - id: hello-world
  query: >
    MATCH
      (m:Message {Message: 'Hello World'})
    RETURN
      m.MessageId AS MessageId,
      m.From AS MessageFrom
```

```
apiVersion: v1
kind: ContinuousQuery
name: message-count
spec:
  mode: query
  sources:
    subscriptions:
      - id: hello-world
  query: >
    MATCH
      (m:Message)
    RETURN
      m.Message AS Message,
      count(m.Message) AS Frequency
```

```
apiVersion: v1
kind: ContinuousQuery
name: inactive-people
spec:
  mode: query
  sources:
    subscriptions:
      - id: hello-world
  query: >
    MATCH
      (m:Message)
    WITH
      m.From AS MessageFrom,
      max(drasi.changeDateTime(m)) AS LastMessageTimestamp
    WHERE
      LastMessageTimestamp <= datetime.realtime() - duration({ seconds: 20 })
    OR
      drasi.trueLater(LastMessageTimestamp <= datetime.realtime() - duration({ seconds:
20 })), LastMessageTimestamp + duration({ seconds: 20 }))
    RETURN
      MessageFrom,
      LastMessageTimestamp
```

```
alan [ ~ ]$ vi hello-world-queries.yaml
alan [ ~ ]$ ./drasi apply -f ./hello-world-queries.yaml
✓ Apply: ContinuousQuery/hello-world-from: complete
✓ Apply: ContinuousQuery/message-count: complete
✓ Apply: ContinuousQuery/inactive-people: complete
```

```
alan [ ~ ]$ ./drasi list query
```

ID	CONTAINER	ERRORMESSAGE	HOSTNAME	STATUS
hello-world-from	default		default-query-host-6959d9446-nnhwf	Running
message-count	default		default-query-host-6959d9446-nnhwf	Running
inactive-people	default		default-query-host-6959d9446-nnhwf	Running

```
alan [ ~ ]$
```


Reaction

```
apiVersion: v1
kind: Reaction
name: hello-world-debug
spec:
  kind: Debug
  queries:
    hello-world-from:
    message-count:
    inactive-people:
```

```
alan [ ~ ]$ ./drasi apply -f ./hello-world-reaction.yaml
✓ Apply: Reaction/hello-world-debug: complete
alan [ ~ ]$ ./drasi list reaction
      ID          | AVAILABLE | MESSAGES
-----+-----+-----
  hello-world-debug | true     |
alan [ ~ ]$
```

Debug

```
kubectl port-forward services/hello-  
world-debug-gateway 8080:8080 -n  
drasi-system
```

debug-reactor

Home

inactive-people

message-count

hello-world-from

[About](#)

Tools Performance Event Stream

Get Debug Info

Query Id	Host
inactive-people	default-query-host-6959d9446-nnhwf
message-count	default-query-host-6959d9446-nnhwf
hello-world-from	default-query-host-6959d9446-nnhwf

Insert Hello World

Which people have sent the message “Hello World”?
How many times has the same message been sent?

```
INSERT INTO public."Message"  
VALUES (5, 'Allen', 'Hello World');
```

```
INSERT INTO public."Message"  
VALUES (5, 'Allen', 'Hello World');
```

debug-reactor

[Home](#)
[inactive-people](#)
[message-count](#)
[hello-world-from](#)

[About](#)

Query Results - hello-world-from

MessageFrom	MessageId
Brian Kernighan	2
Allen	5
Allen2	6

Debug Info

hostName: default-query-host-6959d9446-nnhwf

status: Running

container: default

errorMessage:

Refetch Result Cache

debug-reactor

[Home](#)
[inactive-people](#)
[message-count](#)
[hello-world-from](#)

[About](#)

Query Results - message-count

Frequency	Message
1	To infinity and beyond!
2	I am Spartacus
3	Hello World

Debug Info

hostName: default-query-host-6959d9446-nnhwf

status: Running

container: default

errorMessage:

Refetch Result Cache

20 seconds inactive

debug-reactor

Home

inactive-people

message-count

hello-world-from

[About](#)

Query Results - inactive-people

LastMessageTimestamp	MessageFrom
1970-01-01 00:00:00 +00:00	Buzz Lightyear
1970-01-01 00:00:00 +00:00	Brian Kernighan
1970-01-01 00:00:00 +00:00	Antoninus
1970-01-01 00:00:00 +00:00	David
2024-12-15 03:40:10.867 +00:00	Allen
2024-12-15 03:40:30.672 +00:00	Allen2

Debug Info

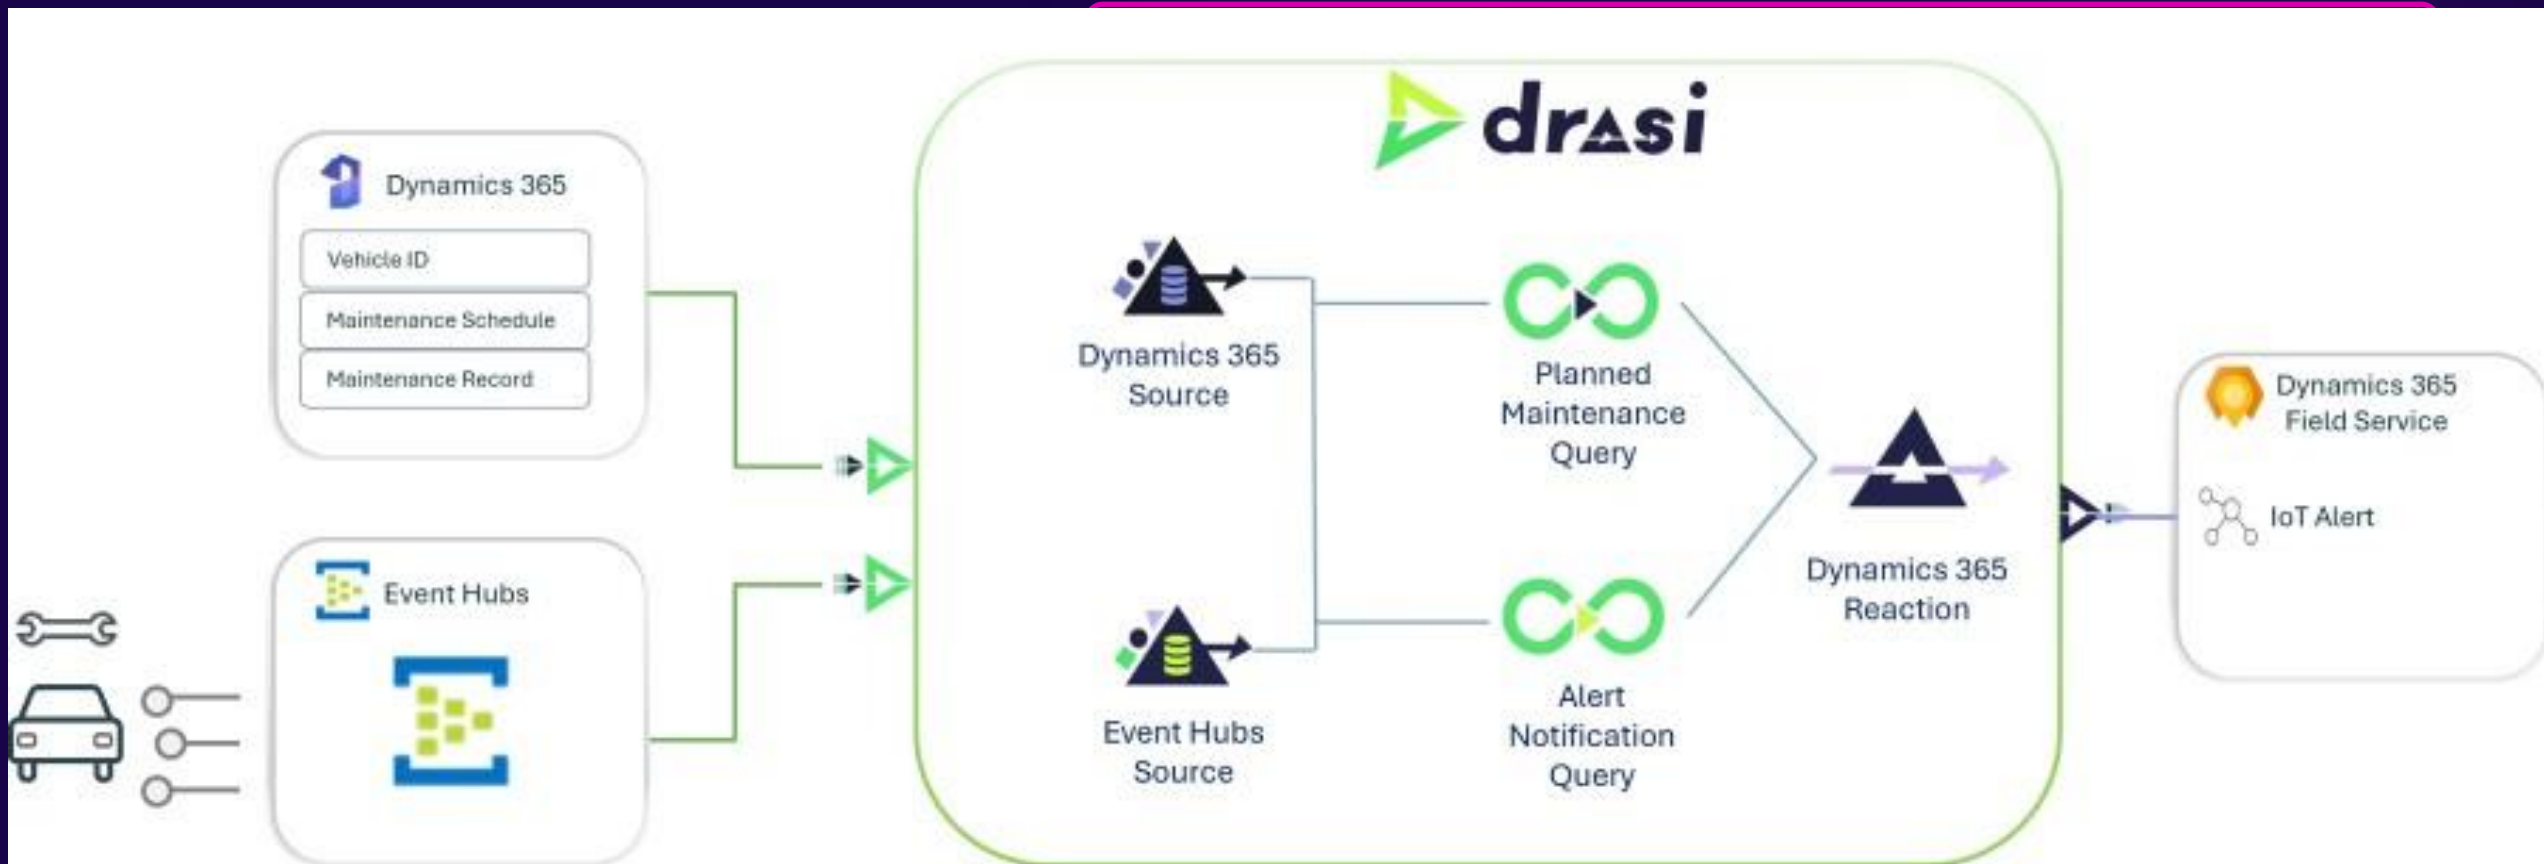
hostName: default-query-host-6959d9446-nnhwf

status: Running

container: default

Conclustion

connected fleet



Solution Design

How to Design Solutions with Drasi

Observing Changes

Entity CRUD

Observing Conditions

an Incident becomes critical.

an Order becomes ready for pickup.

a Room's temperature exceeds 80 degrees.

the occupancy of a Store exceeds 100 people.

Observing Collections

There are new customer orders. These need to be picked from stock.

Once picked, orders need to be packed and prepared for dispatch.

Once prepared, the orders need to be dispatched through various delivery options.

Once dispatched, the orders need to be tracked until delivery.

Once the Order is delivered, it is complete, unless the customer has an issue, in which case a customer support process is initiated...

Solution for Down Stream Application

What is it not good for

01

has a mature change
notification capability

02

Continuous Query
includes data types
for which there is a
lot of data

03

stream analytics or
streaming data
transformation over
high volume data
streams

Reference


1. [Drasi: Microsoft's newest open-source project simplifies change detection and reaction in complex systems](#)
2. [Detect and react intelligently to changes in data with Drasi](#)
3. [Drasi.io](#)
4. [Learning](#)



STUDY4

為 學 習 而 生

特別感謝



以及各位參與活動的各位

